

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

# 14. DataSet Iris ¶

In [2]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\14_Iris.csv")
a
```

Out[2]:

	<b>Id</b>	<b>SepalLengthCm</b>	<b>SepalWidthCm</b>	<b>PetalLengthCm</b>	<b>PetalWidthCm</b>	<b>Species</b>
<b>0</b>	1	5.1	3.5	1.4	0.2	Iris-setosa
<b>1</b>	2	4.9	3.0	1.4	0.2	Iris-setosa
<b>2</b>	3	4.7	3.2	1.3	0.2	Iris-setosa
<b>3</b>	4	4.6	3.1	1.5	0.2	Iris-setosa
<b>4</b>	5	5.0	3.6	1.4	0.2	Iris-setosa
...	...	...	...	...	...	...
<b>145</b>	146	6.7	3.0	5.2	2.3	Iris-virginica
<b>146</b>	147	6.3	2.5	5.0	1.9	Iris-virginica
<b>147</b>	148	6.5	3.0	5.2	2.0	Iris-virginica
<b>148</b>	149	6.2	3.4	5.4	2.3	Iris-virginica
<b>149</b>	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

In [3]:

```
b=a.head(10)
b
```

Out[3]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
5	6	5.4	3.9	1.7	0.4	Iris-setosa
6	7	4.6	3.4	1.4	0.3	Iris-setosa
7	8	5.0	3.4	1.5	0.2	Iris-setosa
8	9	4.4	2.9	1.4	0.2	Iris-setosa
9	10	4.9	3.1	1.5	0.1	Iris-setosa

In [4]:

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id              150 non-null   int64
1   SepalLengthCm   150 non-null   float64
2   SepalWidthCm    150 non-null   float64
3   PetalLengthCm   150 non-null   float64
4   PetalWidthCm    150 non-null   float64
5   Species         150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

In [5]:

```
a.describe()
```

Out[5]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
<b>count</b>	150.000000	150.000000	150.000000	150.000000	150.000000
<b>mean</b>	75.500000	5.843333	3.054000	3.758667	1.198667
<b>std</b>	43.445368	0.828066	0.433594	1.764420	0.763161
<b>min</b>	1.000000	4.300000	2.000000	1.000000	0.100000
<b>25%</b>	38.250000	5.100000	2.800000	1.600000	0.300000
<b>50%</b>	75.500000	5.800000	3.000000	4.350000	1.300000
<b>75%</b>	112.750000	6.400000	3.300000	5.100000	1.800000
<b>max</b>	150.000000	7.900000	4.400000	6.900000	2.500000

In [6]:

```
a.columns
```

Out[6]:

```
Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
      'Species'],  
      dtype='object')
```

In [7]:

```
c=b[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]  
c
```

Out[7]:

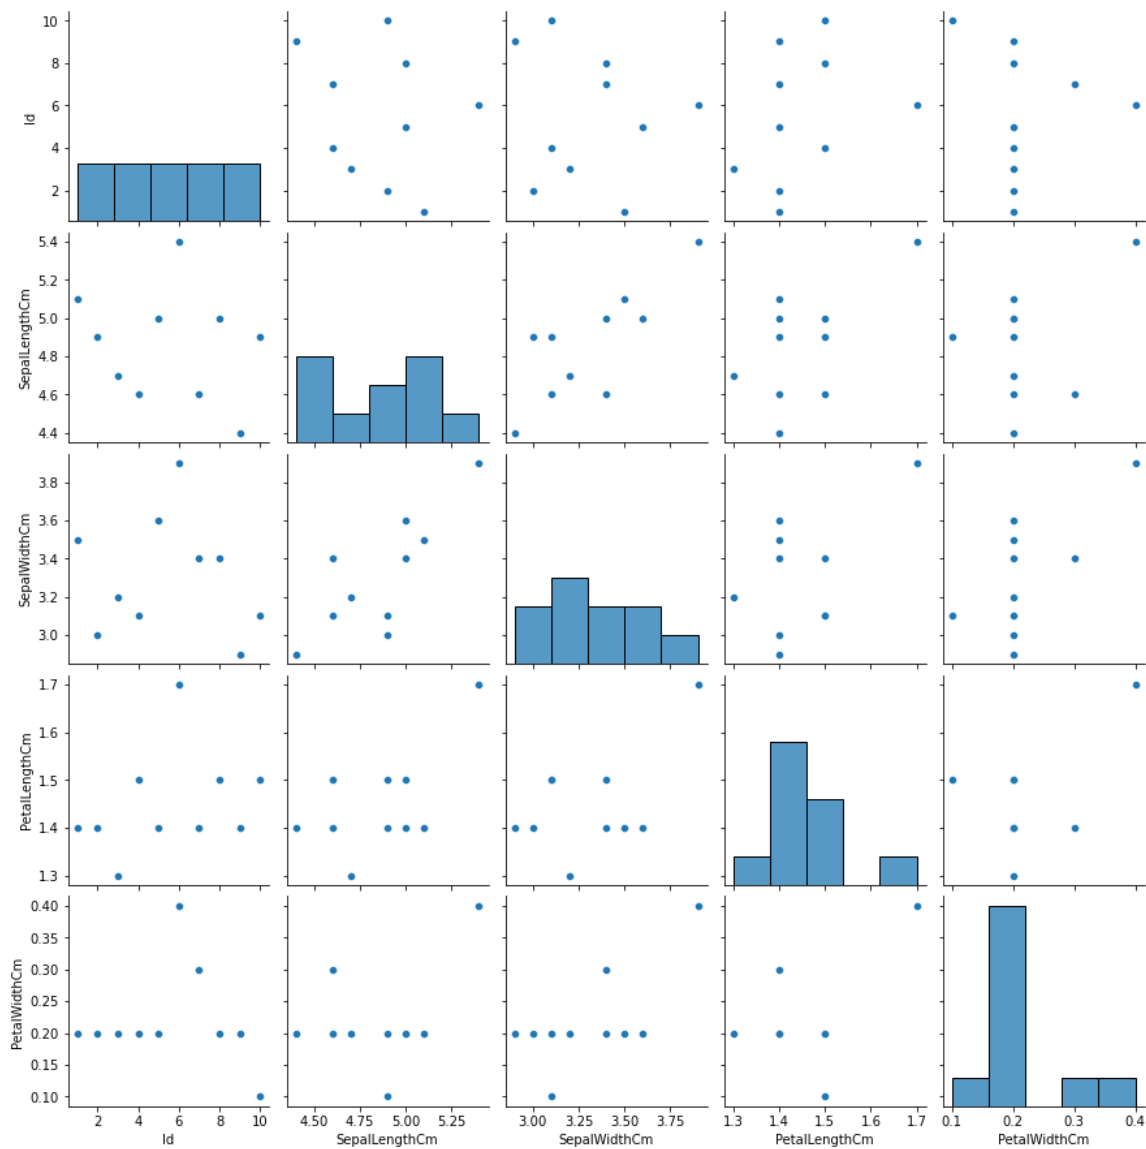
	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
<b>0</b>	1	5.1	3.5	1.4	0.2
<b>1</b>	2	4.9	3.0	1.4	0.2
<b>2</b>	3	4.7	3.2	1.3	0.2
<b>3</b>	4	4.6	3.1	1.5	0.2
<b>4</b>	5	5.0	3.6	1.4	0.2
<b>5</b>	6	5.4	3.9	1.7	0.4
<b>6</b>	7	4.6	3.4	1.4	0.3
<b>7</b>	8	5.0	3.4	1.5	0.2
<b>8</b>	9	4.4	2.9	1.4	0.2
<b>9</b>	10	4.9	3.1	1.5	0.1

In [8]:

```
sns.pairplot(c)
```

Out[8]:

&lt;seaborn.axisgrid.PairGrid at 0x217d2587ac0&gt;



In [9]:


```
sns.distplot(c['SepalLengthCm'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

```
<AxesSubplot:xlabel='SepalLengthCm', ylabel='Density'>
```




In [10]:

```
sns.heatmap(c.corr())
```

Out[10]:

```
<AxesSubplot:>
```



In [11]:

```
x=b[['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]  
y=b['SepalWidthCm']
```

In [12]:

```
from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [13]:

```
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[13]:

LinearRegression()

In [14]:

```
print(lr.intercept_)  
  
-8.881784197001252e-16
```

In [15]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[15]:


	Co-efficient
Id	0.000000e+00
SepalLengthCm	-4.724091e-18
SepalWidthCm	1.000000e+00
PetalLengthCm	-3.819093e-16
PetalWidthCm	1.138951e-16

In [16]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[16]:

<matplotlib.collections.PathCollection at 0x217d45a1c40>



In [17]:

```
print(lr.score(x_test,y_test))
```

1.0

In [18]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [19]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[19]:

Ridge(alpha=10)

In [20]:

```
rr.score(x_test,y_test)
```

Out[20]:

-9.430062712045332

In [21]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[21]:

Lasso(alpha=10)

In [22]:

```
la.score(x_test,y_test)
```

Out[22]:

-10.042274052478106

In [23]:

```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[23]:

ElasticNet()

In [24]:

```
print(en.coef_)
```

[-0. 0. 0. 0. 0.]

In [25]:

```
print(en.intercept_)
```

3.4285714285714284

In [26]:

```
prediction=en.predict(x_test)
print(en.score(x_test,y_test))
```

-10.042274052478106

## Evaluation Metrics

In [27]:

```
from sklearn import metrics
```

In [28]:

```
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 0.395238095238095

In [29]:

```
print("mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

mean Squared Error: 0.17176870748299308



In [30]:

```
print("Root mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root mean Squared Error: 0.414449885369743

## 15. DataSet Horse\_Racing


In [31]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\15_Horse Racing Results.CSV - 15_Horse Racing Res  
a
```

Out[31]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	C
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	
...	...	...	...	...	...	...	...	...	...	
27003	14.06.2020	Sha Tin	11	1200	Gress	1450000	6	A Hamelin	59	A
27004	21.06.2020	Sha Tin	2	1200	Gress	967000	7	K C Leung	57	A
27005	21.06.2020	Sha Tin	4	1200	Gress	967000	6	Blake Shinn	57	A
27006	21.06.2020	Sha Tin	5	1200	Gress	967000	14	Joao Moreira	57	2
27007	21.06.2020	Sha Tin	11	1200	Gress	1450000	7	C Schofield	55	2

27008 rows × 21 columns



In [32]:

```
b=a.head(10)
b
```

Out[32]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Count
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sverig
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sverig
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sverig
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sverig
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sverig
5	10.12.2017	Sha Tin	1	1800	Gress	1310000	4	C Y Ho	52	Sverig
6	01.01.2018	Sha Tin	9	1800	Gress	1310000	9	C Schofield	54	Sverig
7	04.02.2018	Sha Tin	5	1800	Gress	1310000	6	Joao Moreira	57	Sverig
8	03.03.2018	Sha Tin	8	1800	Gress	1310000	3	C Y Ho	56	Sverig
9	11.03.2018	Sha Tin	10	1600	Gress	1310000	8	C Y Ho	57	Sverig

10 rows × 21 columns



In [33]:

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27008 entries, 0 to 27007
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Dato                   27008 non-null  object
 1   Track                  27008 non-null  object
 2   Race Number           27008 non-null  int64
 3   Distance               27008 non-null  int64
 4   Surface                27008 non-null  object
 5   Prize money            27008 non-null  int64
 6   Starting position      27008 non-null  int64
 7   Jockey                 27008 non-null  object
 8   Jockey weight          27008 non-null  int64
 9   Country                27008 non-null  object
10   Horse age              27008 non-null  int64
11   TrainerName            27008 non-null  object
12   Race time              27008 non-null  object
13   Path                   27008 non-null  int64
14   Final place            27008 non-null  int64
15   FGrating               27008 non-null  int64
16   Odds                   27008 non-null  object
17   RaceType               27008 non-null  object
18   HorseId                27008 non-null  int64
19   JockeyId               27008 non-null  int64
20   TrainerID              27008 non-null  int64
dtypes: int64(12), object(9)
memory usage: 4.3+ MB
```

In [34]:

```
a.describe()
```

Out[34]:

	Race Number	Distance	Prize money	Starting position	Jockey weight	Horse age
count	27008.000000	27008.000000	2.700800e+04	27008.000000	27008.000000	27008.000000
mean	5.268624	1401.666173	1.479445e+06	6.741447	55.867373	5.246408
std	2.780088	276.065045	2.162109e+06	3.691071	2.737006	1.519880
min	1.000000	1000.000000	6.600000e+05	1.000000	47.000000	2.000000
25%	3.000000	1200.000000	9.200000e+05	4.000000	54.000000	4.000000
50%	5.000000	1400.000000	9.670000e+05	7.000000	56.000000	5.000000
75%	8.000000	1650.000000	1.450000e+06	10.000000	58.000000	6.000000
max	11.000000	2400.000000	2.800000e+07	14.000000	63.000000	12.000000

In [35]:

```
a.columns
```

Out[35]:

```
Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
      'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
      'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',
      'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],
      dtype='object')
```

In [36]:

```
c=b[['Race Number', 'Distance', 'Surface', 'Prize money',
      'Starting position']]
c
```

Out[36]:

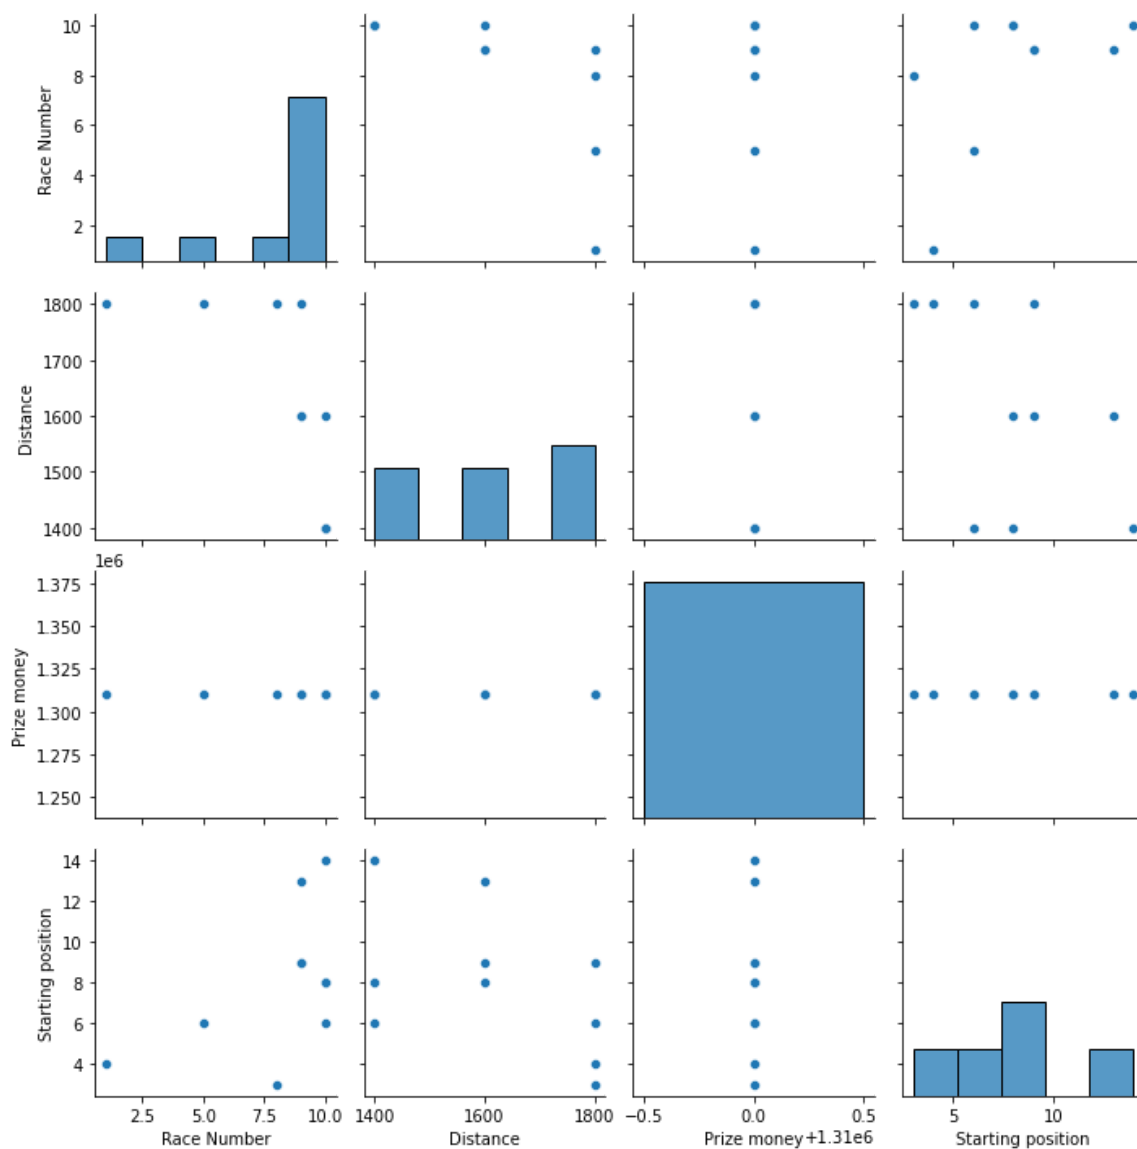
	Race Number	Distance	Surface	Prize money	Starting position
0	10	1400	Gress	1310000	6
1	10	1400	Gress	1310000	14
2	10	1400	Gress	1310000	8
3	9	1600	Gress	1310000	13
4	9	1600	Gress	1310000	9
5	1	1800	Gress	1310000	4
6	9	1800	Gress	1310000	9
7	5	1800	Gress	1310000	6
8	8	1800	Gress	1310000	3
9	10	1600	Gress	1310000	8

In [37]:

```
sns.pairplot(c)
```

Out[37]:

&lt;seaborn.axisgrid.PairGrid at 0x217d4a2c160&gt;



In [38]:


```
sns.distplot(c['Distance'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[38]:

```
<AxesSubplot:xlabel='Distance', ylabel='Density'>
```



In [39]:

```
sns.heatmap(c.corr())
```

Out[39]:

```
<AxesSubplot:>
```



In [40]:

```
x=b[['Race Number', 'Distance', 'Prize money',  
     'Starting position']]  
y=b['Path']
```

In [41]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [42]:

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[42]:

LinearRegression()

In [43]:

```
print(lr.intercept_)
```

-16.134030197444815

In [44]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[44]:


	Co-efficient
<b>Race Number</b>	0.660163
<b>Distance</b>	0.006267
<b>Prize money</b>	0.000000
<b>Starting position</b>	0.240418

In [45]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[45]:

<matplotlib.collections.PathCollection at 0x217d57e49a0>



In [46]:

```
print(lr.score(x_test,y_test))
```

-17.078783064286586

In [47]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [48]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[48]:

Ridge(alpha=10)

In [49]:

```
rr.score(x_test,y_test)
```

Out[49]:

-7.724029079501507

In [50]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[50]:

Lasso(alpha=10)



In [51]:

```
la.score(x_test,y_test)
```

Out[51]:

-0.8690625000000003

In [52]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[52]:

ElasticNet()

In [53]:

```
print(en.coef_)
```

[0.24550125 0.00280497 0. 0.18344861]

In [54]:

```
print(en.intercept_)
```

-6.434565790561453

In [55]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

-5.48654541826254

## Evaluation Metrics

In [56]:

```
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 1.8230175459888072

In [57]:

```
print("mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

mean Squared Error: 4.324363612175026

In [58]:

```
print("Root mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root mean Squared Error: 2.0795104260799047

## 16. DataSet Sleep\_Health


In [59]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\16_Sleep_health_and_lifestyle_dataset.csv")
a
```

Out[59]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pr
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	
...	...	...	...	...	...	...	...	...	...	...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	

374 rows × 13 columns



In [60]:

```
b=a.head(10)
b
```

Out[60]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	B Pres
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	12
1	2	Male	28	Doctor	6.2	6	60	8	Normal	12
2	3	Male	28	Doctor	6.2	6	60	8	Normal	12
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	14
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	14
5	6	Male	28	Software Engineer	5.9	4	30	8	Obese	14
6	7	Male	29	Teacher	6.3	6	40	7	Obese	14
7	8	Male	29	Doctor	7.8	7	75	6	Normal	12
8	9	Male	29	Doctor	7.8	7	75	6	Normal	12
9	10	Male	29	Doctor	7.8	7	75	6	Normal	12

In [61]:

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Person ID                            374 non-null    int64
1   Gender                               374 non-null    object
2   Age                                   374 non-null    int64
3   Occupation                           374 non-null    object
4   Sleep Duration                       374 non-null    float64
5   Quality of Sleep                     374 non-null    int64
6   Physical Activity Level               374 non-null    int64
7   Stress Level                         374 non-null    int64
8   BMI Category                         374 non-null    object
9   Blood Pressure                       374 non-null    object
10  Heart Rate                           374 non-null    int64
11  Daily Steps                          374 non-null    int64
12  Sleep Disorder                       374 non-null    object
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

In [62]:

a.describe()

Out[62]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate
<b>count</b>	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
<b>mean</b>	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775
<b>std</b>	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676
<b>min</b>	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000
<b>25%</b>	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000
<b>50%</b>	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000
<b>75%</b>	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000
<b>max</b>	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000

In [63]:

a.columns

Out[63]:

```
Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',
      'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
      'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',
      'Sleep Disorder'],
      dtype='object')
```

In [64]:

```
c=b[['Sleep Duration',
      'Quality of Sleep', 'Physical Activity Level', 'Stress Level']]
c
```

Out[64]:


	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level
<b>0</b>	6.1	6	42	6
<b>1</b>	6.2	6	60	8
<b>2</b>	6.2	6	60	8
<b>3</b>	5.9	4	30	8
<b>4</b>	5.9	4	30	8
<b>5</b>	5.9	4	30	8
<b>6</b>	6.3	6	40	7
<b>7</b>	7.8	7	75	6
<b>8</b>	7.8	7	75	6
<b>9</b>	7.8	7	75	6

In [65]:

```
sns.pairplot(c)
```

Out[65]:

&lt;seaborn.axisgrid.PairGrid at 0x217d583d400&gt;



In [66]:


```
sns.distplot(b['Daily Steps'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[66]:

<AxesSubplot:xlabel='Daily Steps', ylabel='Density'>

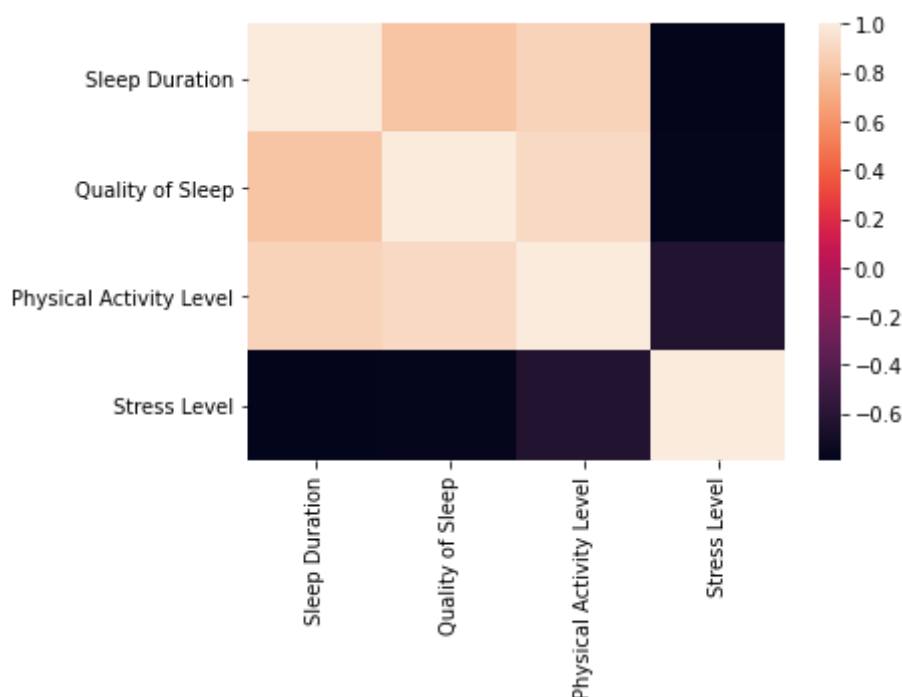


In [67]:

```
sns.heatmap(c.corr())
```

Out[67]:

<AxesSubplot:>



In [68]:

```
x=b[['Age', 'Sleep Duration',  
     'Quality of Sleep', 'Physical Activity Level', 'Stress Level']]  
y=b['Heart Rate']
```

In [69]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [70]:

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[70]:

LinearRegression()

In [71]:

```
print(lr.intercept_)
```

92.64886316724176

In [72]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[72]:


	Co-efficient
<b>Age</b>	0.085580
<b>Sleep Duration</b>	-0.068389
<b>Quality of Sleep</b>	0.196109
<b>Physical Activity Level</b>	-0.345723
<b>Stress Level</b>	-0.006792

In [73]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[73]:

<matplotlib.collections.PathCollection at 0x217d3cc96d0>



In [74]:

```
print(lr.score(x_test,y_test))
```

0.5948164949109942

In [75]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [76]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[76]:

Ridge(alpha=10)

In [77]:

```
rr.score(x_test,y_test)
```

Out[77]:

0.6107776395910929

In [78]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[78]:

Lasso(alpha=10)



In [79]:

```
la.score(x_test,y_test)
```

Out[79]:

0.676799228718691

In [80]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[80]:

ElasticNet()

In [81]:

```
print(en.coef_)
```

```
[ 0.          -0.          0.          -0.33422547 -0.          ]
```

In [82]:

```
print(en.intercept_)
```

95.100509964194

In [83]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.6128318822957425

## Evaluation Metrics

In [84]:

```
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 1.403510675491584

In [85]:

```
print("mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

mean Squared Error: 5.50639100734944

In [86]:

```
print("Root mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root mean Squared Error: 2.3465700516603887

## 17. DataSet Student\_Marks

In [87]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\17_student_marks.csv")  
a
```

Out[87]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
0	22000	78	87	91	91	88	98	94	100	100	100
1	22001	79	71	81	72	73	68	59	69	59	60
2	22002	66	65	70	74	78	86	87	96	88	88
3	22003	60	58	54	61	54	57	64	62	72	60
4	22004	99	95	96	93	97	89	92	98	91	91
5	22005	41	36	35	28	35	36	27	26	19	20
6	22006	47	50	47	57	62	64	71	75	85	88
7	22007	84	74	70	68	58	59	56	56	64	70
8	22008	74	64	58	57	53	51	47	45	42	40
9	22009	87	81	73	74	71	63	53	45	39	40
10	22010	40	34	37	33	31	35	39	38	40	40
11	22011	91	84	78	74	76	80	80	73	75	70
12	22012	81	83	93	88	89	90	99	99	95	88
13	22013	52	50	42	38	33	30	28	22	12	20
14	22014	63	67	65	74	80	86	95	96	92	88
15	22015	76	82	88	94	85	76	70	60	50	50
16	22016	83	78	71	71	77	72	66	75	66	60
17	22017	55	45	43	38	43	35	44	37	45	50
18	22018	71	67	76	74	64	61	57	64	61	50
19	22019	62	61	53	49	54	59	68	74	65	50
20	22020	44	38	36	34	26	34	39	44	36	40
21	22021	50	56	53	46	41	38	47	39	44	50
22	22022	57	48	40	45	43	36	26	19	9	10
23	22023	59	56	52	44	50	40	45	46	54	50
24	22024	84	92	89	80	90	80	84	74	68	70
25	22025	74	80	86	87	90	100	95	87	85	70
26	22026	92	84	74	83	93	83	75	82	81	70
27	22027	63	70	74	65	64	55	61	58	48	40
28	22028	78	77	69	76	78	74	67	69	78	60
29	22029	55	58	59	67	71	62	53	61	67	70
30	22030	54	54	48	38	35	45	46	47	41	50
31	22031	84	93	97	89	86	95	100	100	100	90
32	22032	95	100	94	100	98	99	100	90	80	88
33	22033	64	61	63	73	63	68	64	58	50	50
34	22034	76	79	73	77	83	86	95	89	90	90
35	22035	78	71	61	55	54	48	41	32	41	40
36	22036	95	89	91	84	89	94	85	91	100	100

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
37	22037	99	89	79	87	87	81	82	74	64	4
38	22038	82	83	85	86	89	80	88	95	87	9
39	22039	65	56	64	62	58	51	61	68	70	7
40	22040	100	93	92	86	84	76	82	74	79	7
41	22041	78	72	73	79	81	73	71	77	83	9
42	22042	98	100	100	93	94	92	100	100	98	9
43	22043	58	62	67	77	71	63	64	73	83	7
44	22044	96	92	94	100	99	95	98	92	84	8
45	22045	86	87	85	84	85	91	86	82	85	8
46	22046	48	55	46	40	34	29	37	34	39	4
47	22047	56	52	54	47	40	35	43	44	40	5
48	22048	42	44	46	53	62	59	57	53	43	5
49	22049	64	54	49	59	54	55	57	59	63	7
50	22050	50	44	37	29	37	46	53	57	55	6
51	22051	70	60	70	62	67	67	68	67	72	6
52	22052	63	73	70	63	60	67	61	59	52	5
53	22053	92	100	100	100	100	100	92	87	94	10
54	22054	64	55	54	61	63	57	47	37	44	4
55	22055	60	66	68	58	49	47	39	29	39	4

In [88]:

```
b=a.head(10)
b
```

Out[88]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6	Test_7	Test_8	Test_9	Test_10
0	22000	78	87	91	91	88	98	94	100	100	100
1	22001	79	71	81	72	73	68	59	69	59	60
2	22002	66	65	70	74	78	86	87	96	88	80
3	22003	60	58	54	61	54	57	64	62	72	60
4	22004	99	95	96	93	97	89	92	98	91	90
5	22005	41	36	35	28	35	36	27	26	19	20
6	22006	47	50	47	57	62	64	71	75	85	80
7	22007	84	74	70	68	58	59	56	56	64	70
8	22008	74	64	58	57	53	51	47	45	42	40
9	22009	87	81	73	74	71	63	53	45	39	40

In [89]:

a.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56 entries, 0 to 55
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Student_ID  56 non-null    int64
 1   Test_1      56 non-null    int64
 2   Test_2      56 non-null    int64
 3   Test_3      56 non-null    int64
 4   Test_4      56 non-null    int64
 5   Test_5      56 non-null    int64
 6   Test_6      56 non-null    int64
 7   Test_7      56 non-null    int64
 8   Test_8      56 non-null    int64
 9   Test_9      56 non-null    int64
10  Test_10     56 non-null    int64
11  Test_11     56 non-null    int64
12  Test_12     56 non-null    int64
dtypes: int64(13)
memory usage: 5.8 KB
```

In [90]:

a.describe()

Out[90]:

	Student_ID	Test_1	Test_2	Test_3	Test_4	Test_5	Test_6
count	56.000000	56.000000	56.000000	56.000000	56.000000	56.000000	56.000000
mean	22027.500000	70.750000	69.196429	68.089286	67.446429	67.303571	66.000000
std	16.309506	17.009356	17.712266	18.838333	19.807179	20.746890	21.054043
min	22000.000000	40.000000	34.000000	35.000000	28.000000	26.000000	29.000000
25%	22013.750000	57.750000	55.750000	53.000000	54.500000	53.750000	50.250000
50%	22027.500000	70.500000	68.500000	70.000000	71.500000	69.000000	65.500000
75%	22041.250000	84.000000	83.250000	85.000000	84.000000	85.250000	83.750000
max	22055.000000	100.000000	100.000000	100.000000	100.000000	100.000000	100.000000

In [91]:

a.columns

Out[91]:


```
Index(['Student_ID', 'Test_1', 'Test_2', 'Test_3', 'Test_4', 'Test_5',
      'Test_6', 'Test_7', 'Test_8', 'Test_9', 'Test_10', 'Test_11',
      'Test_12'],
      dtype='object')
```

In [92]:

```
sns.pairplot(b)
```

Out[92]:

&lt;seaborn.axisgrid.PairGrid at 0x217d65aa6d0&gt;



In [93]:


```
sns.distplot(b['Test_2'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[93]:

```
<AxesSubplot:xlabel='Test_2', ylabel='Density'>
```




In [94]:

```
sns.heatmap(b.corr())
```

Out[94]:

```
<AxesSubplot:>
```



In [95]:

```
x=b[['Student_ID', 'Test_1', 'Test_2', 'Test_3', 'Test_4', 'Test_5',  
     'Test_6', 'Test_7', 'Test_8', 'Test_9', 'Test_10', 'Test_11',  
     'Test_12']]  
y=b['Test_6']
```

In [96]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [97]:

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[97]:

LinearRegression()

In [98]:

```
print(lr.intercept_)
```

13.200685689397119

In [99]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[99]:

	Co-efficient
Student_ID	-0.000308
Test_1	-0.279774
Test_2	0.055882
Test_3	0.299643
Test_4	0.157631
Test_5	-0.019157
Test_6	0.500782
Test_7	0.008546
Test_8	0.139421
Test_9	0.108753
Test_10	0.073561
Test_11	-0.014964
Test_12	-0.128484




In [100]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[100]:

<matplotlib.collections.PathCollection at 0x217df4ccd60>



In [101]:

```
print(lr.score(x_test,y_test))
```

0.9695653774078994

In [102]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [103]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[103]:

Ridge(alpha=10)

In [104]:

```
rr.score(x_test,y_test)
```

Out[104]:

0.9644727349653939

In [105]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[105]:

Lasso(alpha=10)

In [106]:

```
la.score(x_test,y_test)
```

Out[106]:

0.9988004699481617

In [107]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[107]:

ElasticNet()

In [108]:

```
print(en.coef_)
```

```
[ -0.          -0.          0.          0.01747778  0.          0.  
  0.93487531  0.01167405  0.03041345  0.          0.          0.  
 -0.          ]
```

In [109]:

```
print(en.intercept_)
```

0.28096938877563105

In [110]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.9996346626033297

## Evaluation Metrics

In [111]:

```
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 0.35516758709910806

In [112]:

```
print("mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

mean Squared Error: 0.15254865963191833

In [113]:

```
print("Root mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root mean Squared Error: 0.39057478110077504

# 18. DataSet World\_Data


In [114]:

```
a=pd.read_csv(r"C:\Users\user\Downloads\18_world-data-2023.csv")
a
```

Out[114]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Call Co
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	9
1	Albania	105	AL	43.10%	28,748	9,000	11.78	35
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	21
3	Andorra	164	AD	40.00%	468	NaN	7.20	37
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	24
...	...	...	...	...	...	...	...	...
190	Venezuela	32	VE	24.50%	912,050	343,000	17.88	5
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	8
192	Yemen	56	YE	44.60%	527,968	40,000	30.45	96
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	26
194	Zimbabwe	38	ZW	41.90%	390,757	51,000	30.68	26

195 rows × 35 columns



In [115]:

```
b=a.fillna(value=51)
b
```

Out[115]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Call Co
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	9
1	Albania	105	AL	43.10%	28,748	9,000	11.78	35
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	21
3	Andorra	164	AD	40.00%	468	51	7.20	37
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	24
...	...	...	...	...	...	...	...	...
190	Venezuela	32	VE	24.50%	912,050	343,000	17.88	5
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	8
192	Yemen	56	YE	44.60%	527,968	40,000	30.45	96
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	26
194	Zimbabwe	38	ZW	41.90%	390,757	51,000	30.68	26

195 rows × 35 columns




In [116]:

```
c=b.head(10)
c
```

Out[116]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0
3	Andorra	164	AD	40.00%	468	51	7.20	376.0
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0
5	Antigua and Barbuda	223	AG	20.50%	443	0	15.33	1.0
6	Argentina	17	AR	54.30%	2,780,400	105,000	17.02	54.0
7	Armenia	104	AM	58.90%	29,743	49,000	13.99	374.0
8	Australia	3	AU	48.20%	7,741,220	58,000	12.60	61.0
9	Austria	109	AT	32.40%	83,871	21,000	9.70	43.0

10 rows × 35 columns



In [117]:

a.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Data columns (total 35 columns):
 #   Column                                                                 Non-Null Count  Dtype
---  -
 0   Country                                                                195 non-null    object
 1   Density                                                                195 non-null    object
    (P/Km2)
 2   Abbreviation                                                            188 non-null    object
 3   Agricultural Land( %)                                                  188 non-null    object
 4   Land Area(Km2)                                                         194 non-null    object
 5   Armed Forces size                                                      171 non-null    object
 6   Birth Rate                                                             189 non-null    float64
 7   Calling Code                                                           194 non-null    float64
 8   Capital/Major City                                                    192 non-null    object
 9   Co2-Emissions                                                         188 non-null    object
10   CPI                                                                    178 non-null    object
11   CPI Change (%)                                                         179 non-null    object
12   Currency-Code                                                         180 non-null    object
13   Fertility Rate                                                         188 non-null    float64
14   Forested Area (%)                                                     188 non-null    object
15   Gasoline Price                                                         175 non-null    object
16   GDP                                                                    193 non-null    object
17   Gross primary education enrollment (%)                                188 non-null    object
18   Gross tertiary education enrollment (%)                              183 non-null    object
19   Infant mortality                                                       189 non-null    float64
20   Largest city                                                           189 non-null    object
21   Life expectancy                                                        187 non-null    float64
22   Maternal mortality ratio                                              181 non-null    float64
23   Minimum wage                                                          150 non-null    object
24   Official language                                                     194 non-null    object
25   Out of pocket health expenditure                                       188 non-null    object
26   Physicians per thousand                                                188 non-null    float64
27   Population                                                             194 non-null    object
28   Population: Labor force participation (%)                             176 non-null    object
29   Tax revenue (%)                                                        169 non-null    object
30   Total tax rate                                                         183 non-null    object
31   Unemployment rate                                                     176 non-null    object
32   Urban_population                                                      190 non-null    object
33   Latitude                                                              194 non-null    float64
34   Longitude                                                             194 non-null    float64
dtypes: float64(9), object(26)
memory usage: 53.4+ KB

```

In [118]:

a.describe()

Out[118]:

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand
<b>count</b>	189.000000	194.000000	188.000000	189.000000	187.000000	181.000000	188.000000
<b>mean</b>	20.214974	360.546392	2.698138	21.332804	72.279679	160.392265	1.839840
<b>std</b>	9.945774	323.236419	1.282267	19.548058	7.483661	233.502024	1.684261
<b>min</b>	5.900000	1.000000	0.980000	1.400000	52.800000	2.000000	0.010000
<b>25%</b>	11.300000	82.500000	1.705000	6.000000	67.000000	13.000000	0.332500
<b>50%</b>	17.950000	255.500000	2.245000	14.000000	73.200000	53.000000	1.460000
<b>75%</b>	28.750000	506.750000	3.597500	32.700000	77.500000	186.000000	2.935000
<b>max</b>	46.080000	1876.000000	6.910000	84.500000	85.400000	1150.000000	8.420000

In [119]:

c.columns

Out[119]:

```
Index(['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land(
%)',
      'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Cod
e',
      'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)',
      'Currency-Code', 'Fertility Rate', 'Forested Area (%)',
      'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)',
      'Gross tertiary education enrollment (%)', 'Infant mortality',
      'Largest city', 'Life expectancy', 'Maternal mortality ratio',
      'Minimum wage', 'Official language', 'Out of pocket health expendit
ure',
      'Physicians per thousand', 'Population',
      'Population: Labor force participation (%)', 'Tax revenue (%)',
      'Total tax rate', 'Unemployment rate', 'Urban_population', 'Latitud
e',
      'Longitude'],
      dtype='object')
```

In [120]:

```
d=c[['Density\n(P/Km2)', 'Urban_population', 'Latitude',
      'Longitude']]
d
```

Out[120]:


	Density\n(P/Km2)	Urban_population	Latitude	Longitude
0	60	9,797,273	33.939110	67.709953
1	105	1,747,593	41.153332	20.168331
2	18	31,510,100	28.033886	1.659626
3	164	67,873	42.506285	1.521801
4	26	21,061,025	-11.202692	17.873887
5	223	23,800	17.060816	-61.796428
6	17	41,339,571	-38.416097	-63.616672
7	104	1,869,848	40.069099	45.038189
8	3	21,844,756	-25.274398	133.775136
9	109	5,194,416	47.516231	14.550072

In [121]:

```
sns.pairplot(d)
```

Out[121]:

&lt;seaborn.axisgrid.PairGrid at 0x217e09599a0&gt;





In [122]:


```
sns.distplot(d['Latitude'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
FutureWarning: `distplot` is a deprecated function and will be removed in  
a future version. Please adapt your code to use either `displot` (a figure  
-level function with similar flexibility) or `histplot` (an axes-level fun  
ction for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[122]:

<AxesSubplot:xlabel='Latitude', ylabel='Density'>




In [123]:

```
sns.heatmap(d.corr())
```

Out[123]:

<AxesSubplot:>



In [124]:

```
x=d[['Density\n(P/Km2)', 'Latitude',  
    'Longitude']]  
y=d['Latitude']
```

In [125]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [126]:

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[126]:

LinearRegression()

In [127]:

```
print(lr.intercept_)
```

-1.0658141036401503e-14

In [128]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[128]:


	Co-efficient
Density\n(P/Km2)	2.906770e-16
Latitude	1.000000e+00
Longitude	-5.332036e-17

In [129]:

```
prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[129]:

<matplotlib.collections.PathCollection at 0x217e0c2fcd0>



In [130]:

```
print(lr.score(x_test,y_test))
```

1.0

In [131]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [132]:

```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[132]:

Ridge(alpha=10)

In [133]:

```
rr.score(x_test,y_test)
```

Out[133]:

0.9996940751514699

In [134]:

```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[134]:

Lasso(alpha=10)

In [135]:

```
la.score(x_test,y_test)
```

Out[135]:

0.9994946274260796

In [136]:

```
en=ElasticNet()  
en.fit(x_train,y_train)
```

Out[136]:

ElasticNet()

In [137]:

```
print(en.coef_)
```

[9.68190444e-04 9.98012216e-01 0.00000000e+00]

In [138]:

```
print(en.intercept_)
```

-0.02666247441751146

In [139]:

```
prediction=en.predict(x_test)  
print(en.score(x_test,y_test))
```

0.9999500554619002

## Evaluation Metrics

In [140]:

```
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 0.07284664765712161

In [141]:

```
print("mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

mean Squared Error: 0.008879604769335657

In [142]:

```
print("Root mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root mean Squared Error: 0.0942316548158614

In [ ]: